


# Combinatorial Optimization by Neural Reinforcement Learning


By Lei Zhang

# Combinatorial optimization

- Finding an optimal solution from a finite set of possible solutions
  - In operations research, applied math, computer science
  - Feasible solutions are discrete, often evolves graphical structures
  - Many are NP-hard



Travelling Salesman Problem (TSP)



Maximum independent set (MIS)

# Machine learning for combinatorial optimization literature

- Routing problems: TSP, Vehicle routing problem(VRP), etc.
  - Vinyals et al., *Pointer networks*, NIPS 2015
    - TSP, Pointer network with **supervised learning**, near optimal solution for 2D TSP up to 50 nodes
  - Bello et al., *Neural combinatorial optimization with reinforcement learning*, ICLR 2017
    - TSP: pointer network, and **RL** (policy gradient), near optimal up to 100 nodes
  - Nazari et al, *Reinforcement Learning for solving the vehicle routing problem*, NeurIPS 2018
    - VRP: RNN, **policy gradient**
  - Kool et al, *Attention, learn to solve routing problem!* ICLR 2019
    - TSP & VPR, Attention network, **policy gradient with baseline**
- Graph theoretical problems: Maximum independent set, minimum vertex cover, maximum cut, etc.
  - Dai et al., *Learning combinatorial optimization algorithms over **graphs***, NeurIPS 2017
    - Structure2vec, greedy algorithm, **DQN**
  - Li et al., *Combinatorial Optimization with **graph convolutional networks** and guided tree search*, NeurIPS 2018
  - Abe et al., *Solving NP-hard problems on graphs by **reinforcement learning** without domain knowledge*
    - Graph Isomorphism Networks and the **Monte-Carlo Tree Search**
- A survey by Benjio, et al.
  - Machine learning for combinatorial optimization: a methodological tour d'horizon, Nov 2018

# Process Flexibility Design Problem

**Stochastic** Mixed Integer Linear Program

**Goal:** design a flexibility structure to maximize expected profit

$$\max_{f \in F} E_{d \sim D}[P(f, d)]$$

- F: set of feasible solutions -- bipartite graph
- D: demand distribution
- P(f, d): a function to calculate maximum profit -- a linear programming problem

# Applications of process flexibility



Manufacturing



Health care



E-commerce



Delivery logistics




Service operations

# Designing flexibility structure


How to design a structure with 15 arcs?

Resource Demand




Design A

Resource Demand



Design B


Resource Demand



Design C

# The neural reinforcement learning approach

- Generating structures like playing games
  - Starting from zero arcs, at each step play an action (add/remove one arc), until target arc # is reached
- State: adjacency matrix
- Action: the arc to be **added/removed**
- Reward:
  - If done: average profits of **50 simulation instances**
  - Else: 0
- Algorithm: Proximal Policy Gradient (PPO)



# Proximal Policy Gradient

Take the biggest possible improvement step on a policy without stepping so far which leads to performance collapse, by penalizing KL-divergence of  $\pi_{\text{old}}$  and  $\pi_{\text{new}}$  or clipping

---

## Algorithm 1 PPO-Clip

---

- 1: Input: initial policy parameters  $\theta_0$ , initial value function parameters  $\phi_0$
- 2: **for**  $k = 0, 1, 2, \dots$  **do**
- 3:   Collect set of trajectories  $\mathcal{D}_k = \{\tau_i\}$  by running policy  $\pi_k = \pi(\theta_k)$  in the environment.
- 4:   Compute rewards-to-go  $\hat{R}_t$ .
- 5:   Compute advantage estimates,  $\hat{A}_t$  (using any method of advantage estimation) based on the current value function  $V_{\phi_k}$ .
- 6:   Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left( \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)} A^{\pi_{\theta_k}}(s_t, a_t), g(\epsilon, A^{\pi_{\theta_k}}(s_t, a_t)) \right),$$

Clipping function

$$g(\epsilon, A) = \begin{cases} (1 + \epsilon)A & A \geq 0 \\ (1 - \epsilon)A & A < 0. \end{cases}$$

Advantage function estimation

typically via stochastic gradient ascent with Adam.

- 7:   Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left( V_{\phi}(s_t) - \hat{R}_t \right)^2,$$


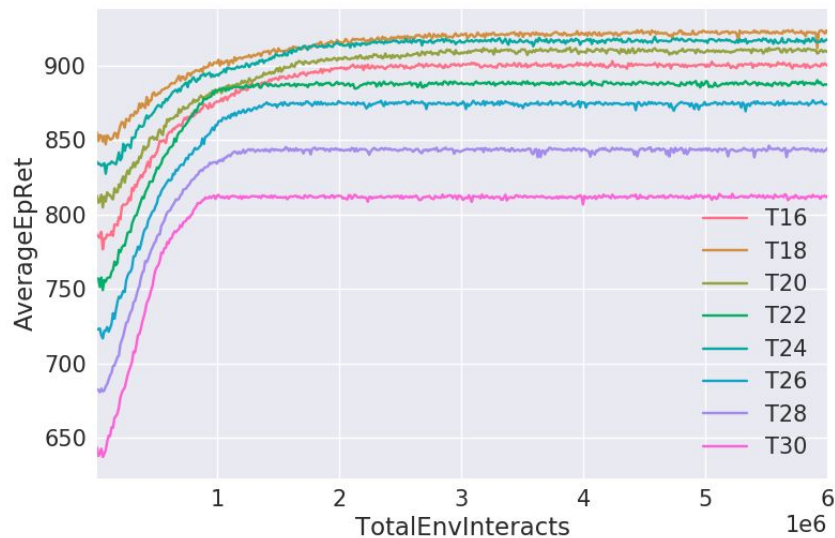
typically via some gradient descent algorithm.

- 8: **end for**
- 

$$\begin{aligned} \hat{A}_t^{(1)} &:= \delta_t^V &= -V(s_t) + r_t + \gamma V(s_{t+1}) \\ \hat{A}_t^{(2)} &:= \delta_t^V + \gamma \delta_{t+1}^V &= -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2}) \\ \hat{A}_t^{(3)} &:= \delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V &= -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 V(s_{t+3}) \\ \hat{A}_t^{(k)} &:= \sum_{l=0}^{k-1} \gamma^l \delta_{t+l}^V &= -V(s_t) + r_t + \gamma r_{t+1} + \dots + \gamma^{k-1} r_{t+k-1} + \gamma^k V(s_{t+k}) \end{aligned}$$




# Training results of 10x10 flexibility environment



# Performance comparison of RL and heuristics

Expected Sales of RL, Drop and Add heuristics



# Main findings

- First attempt to apply neural RL to stochastic combinatorial optimization problems
  - Stochasticity plays to RL's advantage
- Game play
  - More sample efficient, and larger learning capacity
  - Others: greedy, or one shot solution
- Extensible to move variations of the same problem
- Takes a long time to develop and tune parameters
  - Justifiable if the gain is high

# Limitations and next steps

- Model is not extensible
  - Graph neural net
  - RNN
- Can also try Monte Carlo Tree Search